

Exemplar: Complete a join

Activity overview

As a security analyst, you'll often find that you need data from more than one table.

Previously, you learned that a **relational database** is a structured database containing tables that are related to each other.

SQL joins enable you to combine tables that contain a shared column. This is helpful when you need to connect information that appears in different tables.

In this lab activity, you'll use SQL joins to connect separate tables and retrieve needed information.

Get ready to apply what you've learned and **join** some data!

This exemplar is a walkthrough of the previous Qwiklab activity, including detailed instructions and solutions. You may use this exemplar if you were unable to complete the lab and/or you need extra guidance in competing lab tasks. You may also refer to this exemplar to prepare for the graded quiz in this module.

Note: The terms **row** and **record** are used interchangeably.

Scenario

In this scenario, you'll investigate a recent security incident that compromised some machines.

You are responsible for getting the required information from the database for the investigation.

Here's how you'll do this task: **First**, you'll use an inner join to identify which employees are using which machines. **Second**, you'll use left and right joins to find machines that do not belong to any specific user and users who do not have any specific machine assigned to them. **Finally**, you'll use an inner join to list all login attempts made by all employees.

You're ready to join tables in SQL!

Note: In this lab you'll be working with the organization database and the tables it contains.

The lab starts with the organization database in the MariaDB shell that is already open. This means you can start with the tasks as soon as you click the **Start Lab** button.

If you unintentionally exit the organization database in the MariaDB shell, you can reconnect by running the **sudo mysql organization** command.

Task 1. Match employees to their machines

First, you must identify which employees are using which machines. The data is located in the **machines** and **employees** tables.

You must use a SQL inner join to return the records you need based on a connecting column. In the scenario, both tables include the **device_id** column, which you'll use to perform the join.

1. Run the following query to retrieve all records from the **machines** table:

```
1 SELECT *
2 FROM machines;
```

You'll note that this query is not sufficient to perform the join and retrieve the information you need.

2. Complete the query to perform an inner join between the **machines** and **employees** tables on the **device_id** column. Replace **X** and **Y** with this column name:

```
1 SELECT *
2 FROM machines
3 INNER JOIN employees ON machines.X = employees.Y;
```

Note: Placing the **employees** table after **INNER JOIN** makes it the right table.

To complete a join you need to link the joined tables on a common column. In the case of the **employees** and **machines** tables, the **device_id** column is common.

The correct query to solve this step:

```
1 SELECT *
2 FROM machines
3 INNER JOIN employees ON machines.device_id = employees.device_id;
```

How many rows did the inner join return?

Answer: The inner join query returned 185 rows.

Task 2. Return more data

You now must return the information on all machines and the employees who have machines. Next, you must do the reverse and retrieve the information of all employees and any machines that are assigned to them.

To achieve this, you'll complete a left join and a right join on the **employees** and **machines** tables. The results will include all records from one or the other table. You must link these tables using the common **device_id** column.

1. Run the following SQL query to connect the **machines** and **employees** tables through a left join. You must replace the keyword **X** in the query:

```
1 SELECT *
2 FROM machines
3 X JOIN employees ON machines.device_id = employees.device_id;
```

The correct query to solve this step:

```
1 SELECT *
2 FROM machines
3 LEFT JOIN employees ON machines.device_id = employees.device_id;
```

Note: In a left join, all records from the table referenced after **FROM** and before **LEFT JOIN** are included in the result. In this case, all records from the **machines** table are included, regardless of whether they are assigned to an employee or not.

What is the value in the username column for the last record returned?

Answer: The last username returned is NULL.

2. Run the following SQL query to connect the **machines** and **employees** tables through a right join. You must replace the keyword **X** in the query to solve the problem:

```
1 SELECT *
2 FROM machines
3 X JOIN employees ON machines.device_id = employees.device_id;
```

Note: In a right join, all records from the table referenced after **RIGHT JOIN** are included in the result. In this case, all records from the **employees** table are included, regardless of whether they have a machine or not.

The correct query to solve this step:

```
1 SELECT *
2 FROM machines
3 RIGHT JOIN employees ON machines.device_id = employees.device_id;
```

What is the value in the username column for the last record returned?

Answer: The value in the username column for the last record returned is areyes.

Task 3. Retrieve login attempt data

To continue investigating the security incident, you must retrieve the information on all employees who have made login attempts. To achieve this, you'll perform an inner join on the **employees** and **log_in_attempts** tables, linking them on the common **username** column.

- Run the following SQL query to perform an inner join on the **employees** and **log_in_attempts** tables. Replace **X** with the name of the right table. Then replace **Y** and **Z** with the name of the column that connects the two tables:

```
1 SELECT *
2 FROM employees
3 INNER JOIN X ON Y = Z;
```

Note: You must specify the table name with the column name (table.column) when joining the tables.

The correct query to solve this step:

```
1  SELECT *
2  FROM employees
3  INNER JOIN log_in_attempts ON employees.username = log_in_attempts.username;
```

How many records are returned by this inner join?

Answer: There are 200 records returned by the inner join.